

A Deep Learning Approach to Discover Router Firmware Vulnerabilities

Amjad Abu-mahfouz, Saed Alrabaae, Mahmoud Khasawneh, Marton Gergely, Kim-Kwang Raymond Choo

Abstract—An attack on Industrial IoT systems can cause severe damage to connected devices and their owners. Therefore, detecting router firmware vulnerabilities has become a critical issue. However, collecting a dataset of firmware samples is challenging as no open-source datasets are available online. A manual effort was required to verify the states of samples in both the Common Vulnerabilities and Exposures (CVE) and The National Vulnerability Database (NVD) databases as either vulnerable or benign. After verification, 1450 samples were collected. This paper investigates the effectiveness of using convolutional neural networks (CNNs) and computer vision techniques to analyze home router firmware. The collected firmware samples were read as an array of byte strings, divided into sub-arrays based on the image's dimensions, then layered on top of one another to produce the firmware images. The images were divided by manufacturer and used as inputs for various CNN models to test their accuracy. Three statistical filtering algorithms were used on each manufacturer's set to produce multiple versions of each set, totaling 24 datasets across four manufacturers, with six datasets per manufacturer (4 filtered images and two grayscale and RGB images). The image filter algorithms used include local binary pattern (LBP), histogram-oriented gradients (HOG), and Gabor filter used on the LBP and HOG sets. After testing all combinations of the filtered/normal datasets with the CNN training model, the HOG filter was the most accurate, with an average accuracy of 85.81% across all tests and models, with results as high as 97.94% when used with the appropriate CNN model.

Index Terms—Industrial Internet of Things Cybersecurity, Router Firmware Vulnerability, CNN for IIoT

I. INTRODUCTION

HOME routers serve as the primary endpoints of the internet, acting as gatekeepers for all data flowing to and from home networks. However, this critical point of failure is often targeted by malicious agents seeking to exploit vulnerabilities in firmware design or in the network protocols it runs, such as DHCP, DNS, FTP, Telnet, and others. Such attacks can lead to serious consequences, including the bypassing of authentications, the exposure of sensitive user information, the discovery of vulnerabilities in home systems and IoT devices, and the exploitation of these vulnerabilities [1]–[10]. With the rapid expansion of the IoT market for residential and industrial applications, securing firmware has become a necessity that cannot be ignored [11]. While protocols have their own vulnerabilities, poor firmware design adds another layer of insecurity that could be easily fixed. Several studies have emphasized the importance of securing

firmware, with proposed solutions ranging from semantic analysis and combinatorial testing to machine learning techniques and image processing [12]–[19].

The classification of firmware binaries is an important issue that has received the attention of many researchers. In [20], the authors used a statistical data flow analysis approach for vulnerability discovery. They searched for common keywords in the front-end used for user input and then used them as reference points for the back-end binaries, which signified input entry points. They then used targeted flow analysis to detect cases of dangerous uses of user input. They used the SaTC library as a base for taint analysis on 39 firmware samples and achieved results that surpassed KARONTE [21], a popular tool for detecting firmware vulnerabilities. In [22], the authors conducted a similar study that classified firmware binaries by reading the string of binary characters of the file as an image and using convolutional neural networks (CNNs) on the images to detect key features for classification. They used a large dataset of over 20,000 samples coupled with 2 image filters across multiple CNN models. After testing, they reported accuracy ranging from 82.34% to 99.97%. Although our manually collected dataset is smaller and focuses only on home router firmware, we aim for a similar accuracy score while using different filters that produce distinctive patterns across unique model architectures.

This study explores the use of machine learning to classify router firmware as either benign or vulnerable, with the aim of identifying any existing vulnerabilities. A dataset of over 2500 firmware samples was used with six different convolutional neural network (CNN) models. Image filters were also applied and compared to determine the best CNN model and filter combination. Specifically, the approach involves converting firmware into images and using image analysis techniques for classification, and the resulting images had noticeable differences, demonstrating the viability of this method as shown in Figure 1. CNN models were then employed to identify features in the home router firmware images and classify them as either vulnerable or benign. Six neural network models were created and tested on the different filtered sets to find the best combination of filter and model. After many rounds of testing, the most successful combination achieved a validation accuracy range of 85% to 98%, consistently the highest across all manufacturers. It is worth mentioning that this approach can be used to detect anomalies in most systems, provided the data can be converted to images and trained on an optimized model.

Our study makes the following contributions:

- To the best of our knowledge, this is the first effort of applying deep learning to the firmware vulnerability detection problem. We designed a firmware vulnerability classifier using CNN models with leveraging the image processing concepts. The experimental results showed that it is possible to classify router firmware with an accuracy of 89% on average.
- We have investigated the effectiveness of using varying CNN models on images made from router firmware samples, these kinds of tests have already been conducted on similar samples. Using the CNN models in combination with several image filtering algorithms and through testing has concluded that

Corresponding Author: Saed Alrabaae (salrabaae@uaeu.ac.ae)

A. Abu-mahfouz is a Bachelor's student at York University, 4700 Keele St, Toronto, ON M3J 1P3, Canada, e-mail: amjada@my.yorku.ca).

S. Alrabaae is with the Information Systems and Security department, College of IT, United Arab Emirates University, UAE, e-mail: salrabaae@uaeu.ac.ae

M. Khasawneh is with the Computer Engineering Department at Alain University, Abu Dhabi, UAE, e-mail: mahmoud.khasawneh@aau.ac.ae

M. Gergely is with the Information Systems and Security department, College of IT, United Arab Emirates University, UAE, e-mail: mgergely@uaeu.ac.ae

K.-K. R. Choo is with the Department of Information System and Cybersecurity, University of Texas at San Antonio, San Antonio, TX 78249, USA, e-mail: raymond.choo@fulbrightmail.org




Fig. 1. The firmware binaries were transformed into images, resulting in a benign version and a vulnerable version of the Asus router model (RT-AX3000), depicted on the left and right sides, respectively. The CNN would be able to detect some discernible distinctions between these two images.

the applied filters can drastically improve the accuracy while reducing the loss to as little as 0.05% in some cases.

- We have used the LBP, HOG, and Gabor filters on a set of over 2500 samples (after augmentation), the results outperform the research projects of much higher sample sizes.

II. CONVOLUTIONAL NEURAL NETWORKS MODEL

Convolutional neural networks (CNNs) are the type of deep learning algorithms that would be most useful for solving this problem, they are used in a variety of applications dealing with image feature classification such as facial recognition. The main structure of most CNNs begins with an input layer where the image is read as a two-dimensional array of pixels, followed by one or more hidden layers where different weights are applied to the pixel values in order to map the image's features, and ending off with an output layer consisting of fully connected layers that create many-to-many relationships amongst the extracted features and test the effectiveness of the applied weight values [23], [24]. The convolutional layer is the heart of the CNN learning model and is where the bulk of computations occur. The process of convolution consists of applying a filter to the image's pixel values; the filter or "kernel" is usually a 3x3 array that consists of the weight values, it is overlapped with the input image and the dot product between the kernel and the image's pixel values is taken. The kernel acts as a sliding window across the entire image, every time it computes the dot product, the resulting value or "node" is added to the outputted feature map [25]. After the convolution, an activation function [26] is applied to the feature map, this is done to normalize the values and add a degree of nonlinearity to the data, thereby allowing the classifier to learn more complex patterns that cannot be learnt through linear representations.

There are multiple activation functions that alter the feature map's data differently, the simplest of which is the linear activation function which applies no transformations. Other well-known nonlinear activation functions [27] include the sigmoid and hyperbolic tangent (tanh) functions, both of which normalize the feature map's values to different ranges. The sigmoid function's normalization range is [0, 1] and tanh is [-1, 1], and it is documented that they do not perform well when normalizing the values that lie around their range's midpoints and decrease the effectiveness of the optimization function at the output layer [28]. The rectified linear unit function (ReLU) is a popular choice in modern CNNs, it increases the effectiveness of the Stochastic gradient descent (SGD), which needs more detailed values that are not diluted by the normalization of the sigmoid and tanh functions. The ReLU function is simple as it only turns all

| | | |
|-----|-----|-----|
| 3.0 | 3.0 | 3.0 |
| 3.0 | 3.0 | 3.0 |
| 3.0 | 2.0 | 3.0 |

| | | | | |
|---|---|---|---|---|
| 3 | 3 | 2 | 1 | 0 |
| 0 | 0 | 1 | 3 | 1 |
| 3 | 1 | 2 | 2 | 3 |
| 2 | 0 | 0 | 2 | 2 |
| 2 | 0 | 0 | 0 | 1 |

Fig. 2. Max pooling

values that are lower than zero into zero, and keeps everything else the same, ultimately retaining more information and allowing more complex relationships in the data to be learned. As such, this was used as the activation function in this article.

The pooling layer is responsible for decreasing the size of the feature map while retaining the structure and dominant features of the data, it has an important role to play in reducing the size and computational cost of processing the data. During this process, a sliding window is moved across the feature map and the maximum value or the average value of the window is taken, resulting in a smaller and compressed feature map [29]. We chose to use the max pooling in this study because it would retain the most significant values of the feature map, while the average pooling method would dilute the feature map's values as shown in Figure 2. It shows how Max pooling is used to condense the important/dominant data of a feature map, thereby removing noise data that may make it harder for the model to learn. Pooling layers are mainly used to reduce the size of the training data and computational power needed to train the model.

The output layer is comprised of fully connected layer(s), these layers represent the data as a one-dimensional array in contrast to the two-dimensional representation that occurs in the convolutional layers, meaning that the feature maps would have to be flattened before going through the output layer. In the fully connected layers, all nodes are connected to all other nodes in the next fully connected layer, creating many-to-many relationship between the nodes or features and finally the different classes; it allows the model to learn complex nonlinear relationships between the features and the classes [29].

The optimization function is placed at the very end of the CNN output layer and is responsible for training the model by updating the weights until the optimal filters are found. The optimization function that is used in this study is the SGD. It measures the change in weight with regard to the change in error. The goal is to find the minimum of the loss function where the gradient or slope reaches close to zero, which would signify the least number of errors. The weight values are updated based on the learning rate (it controls magnitude of change) and then they are backpropagated through the network and tested again with every epoch allowing the SGD to optimize further [30].

III. IMAGE TYPES AND FILTERS

In the field of computer vision, image texture can be defined as the variation of pixel brightness intensity in terms of space/distance, and is the main way to define objects or concepts in a given image. Texture analysis is part of computer vision, and deals with object or feature recognition, medical image analysis, pattern recognition, etc. Texture analysis has many approaches to detect image features




Fig. 3. Image from Asus's ZenWiFi series

and can be further divided into four sub-categories: statistical methods, structural methods, model-based methods, and transform based methods. The statistical approach [31] divides an image into separate areas and derives various statistics from the surrounding pixel values in each area and would be the best fit for this study. Other methods such as structural analysis which deal with recognizing repeating patterns are not optimized for this.

A. Making RGB and Grayscale Images

To convert the firmware binary files into images, a simple script was made in python. The program reads the bytes of the binary file into a buffer, the buffer is then converted into a one-dimensional unit8 array which consists of the byte values [0, 255]. In order to make a square (equally sized) two-dimensional array from it, the square root of the size of the array is taken and then used in a ceiling function to obtain the width and height of the new 2-D array; the new size is calculated and the extra space is padded with zeros in unit8 format. Finally, the resized and padded 1-D array is reshaped into a 2D through the numpy reshape module and then converted to an image using pillow's image, from array() method.

The model takes in a byte stream and makes a [n x n] dimensional image array; the extra padding at the end is to make sure the 2D image array has a square shape. We turned the firmware to images because from the beginning.

We wanted to test if CNNs could be effective in classifying them as in other applications like facial recognition.

Figure 3 shows one of the images converted from Asus's ZenWiFi series. There are distinct horizontal bands on the top section with less defined diagonal bands the CNN model would easily detect these and others that are not discernable by the human eye.

B. HOG Filter

Histogram oriented gradients specialize in detecting the shape and structure of an image's features. This is in contrast to other feature extraction methods that look for edge information only, which HOG can also find through extracted gradient and orientation information. The gradient information is obtained through dividing the image into patches or areas.

Then looking at the center pixel and the ones surrounding it within the area - in different directions and varying lengths the difference of pixel intensity is computed to get the gradient of the directions. After the gradients have been obtained, the gradient magnitudes are calculated through the Pythagoras theorem while the directions are obtained through the arctan function [32]. Figure 4 shows the HOG filter that is able to condense an image into its local directional components in terms of pixel intensity. It creates an image with distinct and well-defined patches. Giving the grainy firmware images higher definition using this model should make it easier to learn.



Fig. 4. HOG Filter




Fig. 5. The LBP algorithm: How an image's pixel values are converted to a decimal value

C. LBP Filter

Histogram oriented gradients specialize in detecting the shape and structure of an image's features. This is in contrast to other feature extraction methods that look for edge information only, which HOG can also find through extracted gradient and orientation information. The gradient information is obtained through dividing the image into patches or areas.

Then looking at the center pixel and the ones surrounding it within the area - in different directions and varying lengths the difference of pixel intensity is computed to get the gradient of the directions. After the gradients have been obtained, the gradient magnitudes are calculated through the Pythagoras theorem while the directions are obtained through the arctan function [32]. Figure 4 shows the HOG filter that is able to condense an image into its local directional components in terms of pixel intensity. It creates an image with distinct and well-defined patches. Giving the grainy firmware images higher definition using this model should make it easier to learn. The local binary pattern is a widely used algorithm for facial recognition and has been performing well ever since its creation in the 1940's [33]. This algorithm works by dividing the image into areas or patches, usually 3x3 and comparing the center pixel of each area to its surrounding pixels. If the surrounding pixel's intensity value is above or equal to the center pixel's value, then the surrounding pixel value is replaced by 1. Otherwise, it is replaced by 0. At this point, the area or matrix is converted from byte values into binary with the center pixel value removed. The surrounding pixel's binary values are concatenated in order (left to right, top to bottom) and the concatenated value is converted from byte to integer, and placed as the center pixel's new value. This continues across all divided areas of the image, with the final result being a new image containing highlighted features of the original image [34]. The LBP algorithm is explained in Figure 5.

It illustrates how an image's pixel values are compared to the central pixel, then converted to binary, and finally into a decimal value. This process uses a sliding window and repeats for all of the image's patches.

D. Gabor Filter

The Gabor filter is a type of linear filter that is modelled after the human visual cortex, these filters have been shown to excel in detecting features that are localized spatially and frequency wise. We believe that the vulnerabilities in the firmware would exhibit a different pixel pattern than the benign versions, such as having some

kind of sanitization step before passing user inputs. Its boasted ability to discern features based on their spatial location and orientation seems like a good fit for analyzing vulnerable or malicious code in firmware images and will be tested. Gabor filters are represented as sinusoidal signals of a particular frequency and orientation that can detect changes in the image's texture [35].

They are modelled based on the equation below.

The inputs of this equation with a description of their function are listed below: Kernel size: this is the size of the Gabor filter and will be set to a 3x3 matrix to extract finer details.

Wavelength (λ): this controls the wavelength of the sinusoidal component, in other words the width of the filter strips. Sigma (σ): this controls the size of the Gabor envelope, if made too large, it could miss the small features entirely while filtering, for these experiments, it was set to 3. Orientation (θ): this controls the orientation of the Gabor strips and will filter the features in those orientations; it was set to 90° for a horizontal alignment of the filter. Phase offset (ϕ): this defines the symmetry and was set to 0 to get a symmetric filter. Aspect ratio (γ): this sets the ellipticity of the filter and was set to 0.5 to make it a thinner and narrower for more detailed features.

E. The Impact of Filters

The reason that filters were used in the first place was to test if they would have an effect on the predictions of the models and if so, which one would benefit the model's predictions the most. Image filters from statistical texture analysis were chosen because they extract statistical information on surrounding pixels in an area; other methods that focus on extracting repeating pixel patterns would be less useful. The main question to be answered was: "would the statistical data be more useful for training the model than the raw pixels?"

LBP, GABOR, and HOG filters all stem from the statistical analysis approach and each one compresses pixel areas in the image based on each filter's algorithms, resulting in a smaller/compressed image while still retaining relevant statistical information.

The HOG and GABOR filters proved to be the most effective ; when used on its own, the HOG filter achieved a highest overall accuracy, even higher than the RGB images (which had the highest accuracy of the unfiltered images) .

IV. FIRMWARE COLLECTION AND EXPERIMENTAL SETUP

A. Firmware Samples

The first step was to collect the firmware samples. We have manually collected this dataset. The CVE and NVD vulnerability database websites were used to search for router firmware vulnerabilities, afterwards we have visited vendors' websites to search for the vulnerable versions. The vendors' sites would mostly feature only the newest benign versions and we had to use external repositories to acquire the vulnerable versions. The entire process took longer than expected, due to most of the firmware being unavailable on the parent site.

At the end, we have managed to acquire samples for 4 main router companies (Asus, Buffalo, D-Link, and Netgear) since they had most of the entries on the CVE and NVD sites while other router companies had too few and not suitable for a large dataset. We have managed to collect over 2400 samples.

The dataset is divided by manufacturer and further divided by vulnerable and benign samples for said manufacturer. The amounts of samples per category are displayed in Table I.

Ultimately, the dataset was divided by router manufacturer for a total of 4 sets; the D-Link dataset had the highest number of samples,

totaling to 891, followed by Asus with a sample size of 612, then Buffalo with 523, and lastly Netgear with 380 samples. Various router models spanning different versions were collected, upon preliminary testing it was determined that some of them would be removed due to being underrepresented in the sets and ultimately decreasing the validation accuracy. The main router model firmware for Asus include the DSL, FW-RT, and the RT series, the buffalo firmware includes the wex, whr, wmr, etc. The D-link router models include DAP, DIR, DSL, and DSR. The Netgear router models include GC108, GS108, Rxxxx.

B. Preprocessing

First, all of the firmware binaries were converted into two separate sets, one grayscale and one RGB by using the image converter python script. Converting the entirety of the dataset to RGB images was considerably longer than converting to greyscale, taking around three times longer. A python program was made to automatically filter the image files using the HOG, LBP, and Gabor filters. These filters with the exception of Gabor only filter 2D grayscale image arrays and the entire process of filtering the grayscale dataset took days to complete with the HOG filter taking the longest of the three. The end result was 24 datasets in total, with 6 for each router manufacturer. Each of the six datasets represents the different image types and filters (Grayscale, RGB, LBP, HOG, LBP+Gabor, HOG+Gabor). Later, during the testing, we could then increase the validation accuracy by decreasing the overfitting of the training set by augmenting the images.

Thus, creating more samples that differ from the original.

C. Assumption

The possibility that the newest benign versions could potentially be classified as vulnerable in the future was apparent from the beginning of this project. We decided to proceed with the explicit assumption that any firmware that was not listed in the nvd archive was to be classified as benign. The rationale for this was that if the vulnerable firmware contained specific pixel patterns that would map to known vulnerabilities, the corresponding benign versions should be missing said patterns.

Using this approach on the dataset containing firmware of multiple routers for a manufacturer could lead to a benign version of router firmware that contains an undiscovered vulnerability that another benign router version patched and would produce noise.

V. TESTING AND RESULTS

A. Preliminary Experiments and Findings

The first round of tests was done on the grayscale image dataset, the purpose of these tests was to figure out what kind of CNN model was going to work for this type of classification problem, and whether we could get reasonable accuracy from the tests. In addition, we are interested to learn more about CNN's hyperparameters and layers, and how they affect the training and accuracy values for our dataset. The first test featured a basic and recommended CNN model, with one hidden layer consisting of one 2D convolution layer with a filter size of 3x3 and 32 filters along with a ReLU as the activation function followed by a max pooling layer and two dense layers for the output with the recommended Adam optimizer. The model was trained on a binary set of classification data consisting of all the company's firmware images, placed in one folder labeled 'benign' and another named 'vulnerable'. This was done over 30 epochs. The results were underwhelming, yielding a final validation accuracy of 47.427%. It took 4 epochs on average to reach the peak validation accuracy, and

then it had infrequent and small spikes up to 49%. After testing with deeper CNN models having a depth of 2 and 3 convolutional layers and different sets of hyperparameters (kernel sizes, number of filters, and different activation functions), the end result was always the same validation accuracy of 47%. At this point, it seemed like the problem was with the dataset and its distribution. After going through the dataset and removing any router model firmware that only had one or two versions/representations, an updated dataset was made.

This was done to streamline the dataset and remove most of the variability that came from a lone router model's firmware that may have interfered with the learning.

Testing with the updated dataset on the first CNN of depth 1 gave a slightly better validation accuracy of 52%. At this point, it was clear that the problem was with the dataset. Comparing the loss of the training and validation sets, it seemed like the model was learning the training set well with its loss decreasing rapidly, but it still struggled with the validation set due to its loss increasing, while fluctuating. This is a clear sign of overfitting, a situation where the validation dataset samples are underrepresented and do not fully mirror the training set samples in the distribution of firmware models. The ways to fix overfitting include implementing dropout layers mainly after the dense layers, using a lower learning rate, including batch normalization, and altering the datasets. The problem was that some router models had many vulnerable versions of firmware but only one or two benign versions for the same model in which case we put them into the training sets for better feature learning. Instead of outright deleting them from the sets the only other option was to use image augmentation to artificially increase their presence and allow us to represent them in both, the training and validation datasets.

The image augmentation was done through the use of the Keras image preprocessing module called ImageDataGenerator. This module allows images to be resized, rescaled, rotated, shifted and zoomed in or out. For the purposes of this study, we decided to apply a vertical shift with a wraparound, along with a slight zoom-in. This translates the image with a similar structure and features from the firmware, but in different positions with the only concerns stemming from the zoom in which would alter the pixel's values and feature structure to enlarge the features.

Testing the same CNN model of depth 1 as in the previous tests with the augmented dataset yielded a validation accuracy of 64% which was an improvement. At this point, we tried playing with different augmented datasets along with different CNN models and hyperparameters. After observing the results, we took note of the following changes: Testing the CNN model on a categorical dataset consisting of all the companies in one dataset but in different categories (ex. 1D linkvulnerable, 1Dlinkbenign, 1Asusvulnerable, ...) results in lower validation accuracy initially, but has a strong upwards trend that eventually surpasses the binary classification's validation accuracy (where all company's benign are in one set, etc.). It also had more stable loss values. Testing each router company's firmware individually as a binary dataset, resulted in the highest accuracy of all by far. This is when we also discovered that the Netgear dataset was the hardest to learn in some cases, and was the main contributor to the low results in the previous tests.

B. The impact of Parameters

Upon conducting multiple experiments of varying hyperparameters and observing the resultant accuracy, we were able to come to some conclusions on their effectiveness. Before conducting the main tests, we tested the impacts of changing the number of filters, using different optimizers and their parameters, implementing dropouts, and the number of dense units.




Fig. 6. These tests used the same model and manufacturer with the first image having a higher filter parameter set while the second image used the optimal amount specified.

1) Number of Filters: Changing the number of filters had noticeable effects on the validation accuracy. We began by using the recommended filter size of 32 and readjusting it after testing in an attempt to increase the accuracy. The accuracy was higher with models that used fewer filters. Such an increase in accuracy with a decrease of the number of filters continued until the filter number was set too low and the accuracy began to slightly decrease. The highest accuracy came from using 5 filters on the first convolutional layer and with deeper CNN models, doubling the filters at each convolutional layer gave optimal results as shown in Figure 6.

This suggests that the images of firmware do not have much complex features that define them and are relatively easy to learn. As shown in Figure 6, both of these tests used the same model and manufacturer with the first image having a higher filter parameter set. While the second image used the optimal amount specified. The second image's test produced an accuracy 2% higher, while not much, this trend was consistent.

2) Line Art figures: Figures that are composed of only black lines and shapes. These figures should have no shades or half-tones of gray, only black and white.

3) Batch normalization layers: Using batch normalization layers after pooling helped increase the accuracy in most cases, although sometimes resulted in a slight decrease of around 2%; another notable difference was in the fluctuations of accuracy.

Using batch normalization layers, the accuracy fluctuations were lower in magnitude, giving more stability to the learning (see Fig. 7). It was decided that using batch normalization in the final tests was best due to smaller fluctuations and would give a clearer picture of the CNN model's performance. Fig. 7 shows the accuracy results having used no batch normalization, while the second result shows results with batch normalization layers. This is a case when no batch normalization layers proved to be more effective, giving a 3% accuracy increase.

These improvements were not across the board. Using no batch normalization layers gave much bigger fluctuations overall.

ACKNOWLEDGMENT

The preferred spelling of the word "acknowledgment" in American English is without an "e" after the "g." Use the singular heading even if you have many acknowledgments. Avoid expressions such as "One of us (S.B.A.) would like to thank" Instead, write "F.




Fig. 7. The impact of batch normalization




Fig. 9. The impact of dense layer




Fig. 8. The impact of dropout layer

A. Author thanks” In most cases, sponsor and financial support acknowledgments are placed in the unnumbered footnote on the first page, not here.

4) *Dropout layer*: Implementing a dropout layer after the dense layers gave higher accuracy in almost all cases (see Fig. 9).

9). Each of the manufacturers and models required a different dropout percentage to produce the optimal accuracy results, for example D-Link needed a 30% dropout while Asus did best with a 25% dropout. It was decided to drop the dropout layers from the final testing phase due to their variability in performance between the manufacturers and different models.

In future testing with more resources, we would like to revisit implementing the dropouts as they produce higher accuracy, but need much more management and overhead. The SGD optimizer was used for this study due to being a popular choice in most CNNs and adjusting its learning rate parameter was the key to a high accuracy. The learning rate was set to the recommended setting of 0.001, and proved to be the best value. Lower learning rates were used and resulted in lower increases in accuracy per epoch, and a slighter lower accuracy in the end, while the higher learning rates decreased the accuracy altogether (see Fig. 8). This shows that learning the firmware’s defining features requires a lower learning rate with smaller adjustments to the filters. Fig. 8 shows when dropout layers are implemented, the regular model used in the main tests with no dropouts and 0.001 learning rate, and the last image is when a higher learning rate of 0.09 is used. It shows smaller fluctuations but

a lower accuracy than the first, to simplify testing, models with no dropouts were used.

5) *Dense layer*: Much like the number of filters, the number of dense layer units needed calibration for higher accuracy, and began to fall off past the optimal value. Preliminary testing began using the recommended 64 dense layer units. From there, decreasing the dense units resulted in lower accuracy, while increasing them to 128 gave the highest accuracy (see Fig. 9). Implementing additional dense layers resulted in lower accuracy, the best overall performance came from a single dense layer consisting of 128 units. Fig. 9 the first image represents the results from using more dense layer units (256), the second image is from using the optimal number (128) and gave a 5% increase in validation accuracy; the tests using the optimal number of dense units gave better accuracy overall.

C. Experimental Results

We divided the dataset by router manufacturer such that training and validation datasets are binary on each router company, for example one of the training sets contains a Dlink vulnerable and D-link benign folders. The LBP and HOG filters only work with grayscale images, they were applied to the grayscale images to create new sets that would also be tested for their accuracy against the RGB and grayscale sets. The datasets that were tested included the grayscale, RGB, HOG, LBP, LBP+Gabor, and HOG+Gabor; each of these contained the testing and training sets for each router company. Two metrics were used in evaluating the success of the CNN models, they are the loss and accuracy of the models.

The loss function used was the binary cross entropy, it uses the logarithmic loss function which compares each prediction to the actual value. A score is then calculated in the range of [0, 1] with smaller scores representing small differences between the prediction and the actual value and in turn more accurate predictions, while the opposite is true for bigger scores. The accuracy metric is calculated by the ratio between the number of correct predictions over the total number of predictions, taking the true positive and true negative over the true and false positives and negatives. These metrics could be deceptive due to the imbalanced representations in the datasets (for example most of the samples belonging to one class and only a few to another class may cause a false accuracy score). However after the image augmentation the sample representations were evened out and should be equally represented. Therefore, we decided to test the sets on four main CNN models with hidden layers and also two models with only dense layers. The reasoning for this was that the firmware

was originally a onedimensional array of bytes and that converting them to images restructured some of the bytes and in turn the features. After the preliminary test's findings, the models that performed the best had convolutional layers with small window (3x3) sizes and fewer filters, the stochastic gradient descent algorithm with a low learning rate had solid performance along with a single dense layer with 128 units and so these hyper parameters are used in the main tests. Each of the models tested are described below: Model-1: This model consists of a single convolutional layer followed by a max pooling layer and then two dense layers.

The convolutional layer has a filter size of 3x3, 5 filters, and a ReLU activation function, the max pooling layer has a window size of 2x2 and a stride of 2, and the dense layer has 128 units while the final dense layer uses an SGD optimizer with a learning rate of 0.001.

Model-2: This model has two consecutive convolutional layers and a pooling layer. This was to test if higher level features could be extracted to give a better accuracy. Using a higher number of filters on the second convolutional layer than the first gave better results than using less or an equal amount.

The first convolutional layer has 5 filters and the second 10, both have a window size of 3x3. This is followed by a max pooling layer and 2 dense layers as in the first model.

Model-3: This model consists of two convolutional layers but they are not consecutive as in the last model and are followed by max pooling layers. The output layer is the same as the previous two models.

Model-4: This model is three convolutional layers deep following the same blueprint as Model-3 where max pooling is applied after each convolution. The first layer has 5 filters followed by 10 in the second, and then 20 in the third.

Model-5: For the next models, we decided to only use dense layers in a one-dimensional representation of the data, preserving its original form and feature structure. Model-5 consists of only one dense layer of 128 units and the output/classification layer.

Model-6: This model is made up of two dense layers, the first having 128 units and the second 256; during the preliminary testing it was determined that when the proceeding dense layer has more units, a better accuracy score is achieved.

D. Main Testing and Findings

From the tests, it was determined that resizing the images to greater dimensions (500x500 vs 250x250) resulted in higher accuracy likely due to more features and details being retained. The datasets of each company had different attributes and were not perfectly equal in their sample representation.

When building the datasets, the main goal was to collect as many samples as possible while only keeping those that had enough representations to positively contribute to the learning algorithm and removing the underrepresented samples. This process resulted in the datasets having varying sizes and differing in number of router model firmware. The performance of the datasets seemed to be tied to the number of samples but not the variability present in them. The Asus and D-Link datasets had the highest number of samples and both performed better than the Buffalo and Netgear sets in most cases including the HOG filtered images, which produced the highest accuracy.

Conversely, the Netgear set which had the fewest number of samples had the lowest accuracy in the HOG filtered tests.

The Asus and Netgear sets had the fewest router models, and yet the Netgear set underperformed and was too inconsistent to make a conclusion based on the variability of the datasets.

The results were reported in Table II.

Using the grayscale dataset produced underwhelming results ranging from 53

The RGB dataset yielded substantially higher accuracy ranging from 73.10

This is likely due to the RGB images retaining more features from the extra channels and allowing the models to learn more complex patterns. The RGB sets delivered some of the highest accuracy and could only be surpassed by the HOG datasets.

The LBP filtered images were made from the grayscale sets and made them look less grainy with the important features highlighted. Applying this filter increased the accuracy of all the models across the board compared to the grayscale sets.

The LBP set accuracies ranged from 59.18

Model 2 was clearly the best fit for the LBP regardless of manufacturer. This demonstrates that the less defined and chaotic structure of the LBP images (compared to the HOG filtered images) require a stacking of convolutional layers to extract higher level features, in contrast with the HOG image features, which are much more distinct and do not require this kind of processing.

The HOG filtered images exhibited the best feature translation as they use areas with distinct features derived from the surrounding pixel intensities and are easily recognizable even by the human eye. During the tests the HOG images produced the highest accuracies of all the datasets from 68.08% to 97.94% with an average accuracy of 85.81%, rivaling and at times surpassing the accuracy of the RGB sets. The HOG filter performed best when used with the Asus dataset, possibly due to having less variability in terms of router model representation. Training Model 4 is the most accurate across all manufacturers. The increased depth allowed for a more detailed feature extraction which paired well with the less abstract HOG filtered images. The Gabor filter was applied to the grayscale, LBP, and HOG filtered images and the parameters of the Gabor filter were set to filter only the horizontal features since the firmware code was read as a horizontal array of byte strings and layered on one another to make the images and would thereby have neighboring code segments in the horizontal orientation. The Gabor filter ultimately had the negative effect of decreasing the accuracy of all the image types that it was applied to, after this conclusion further testing on it ended; although the Gabor filtered images exhibited a decrease in accuracy from their original counterparts, the ranking in their accuracy was carried over with the Gabor filtered HOG images having the highest followed by Gabor+LBP. When the one-dimensional models 5 and 6 were trained, both performed poorly in comparison to all the previous models that used two dimensional representations of the data, this was across all manufacturers and filters/image types. The best performance in terms of accuracy came from the HOG filtered datasets when used with Model 4, this proves that vulnerable code segments are more easily detectable when using multiple convolutional layers separated by pooling layers as they narrow down the feature search by removing unnecessary noise data. This is mainly true for the HOG filtered images and possibly others not yet tested that have well defined features with little entropy and noise data.

VI. ADDITIONAL EXPERIMENTS

A. classification results can be generalized to other manufacturers or models of routers

This issue was explored in the latest rounds of testing that dealt with the models learning to differentiate between the types of vulnerabilities (regardless of vendor) in contrast to the previous training sessions that only had a vulnerable/benign binary classification. The

testing done featured using a new categorical dataset that would divide the firmware based on types of vulnerabilities; for this test the three major types of vulnerabilities with the most number of samples were selected and used as separate categories [Access control bypass, Arbitrary code execution, and SQL injection], in addition three more categories would be added to represent their benign patched counterparts.

The test concluded with promising results; although it ended with a 78.86% accuracy at epoc 20, that was mainly due to overfitting and faulty hyperparameters. Prior to reaching the 16th epoc, the model was frequently scoring accuracy in the mid 87%'s before it began to drop.

This test proved that the model has potential at learning different types of vulnerabilities regardless of vendor and with more adjustments to the model and better representation in the dataset, the accuracy should achieve a higher, more stable accuracy.

B. Discovering Unknown Vulnerabilities

We tested how well this model predicts firmware that it has never encountered. For this experiment, we decided to use the buffalo dataset since it had the best firmware representation.

While the other vendor datasets had gaps in discontinued or legacy firmware. RGB images used with model 2 provided the best results and were used for this test (the input size and all of the hyperparameters being the same as well).

The latest versions of firmware were removed from the benign and vulnerable training sets and added to the testing sets. This was done in order to simulate the model not encountering the firmware during training. If the firmware version was the only version present in the dataset, it was left alone and if a version was taken out, all of its augmentations would also be removed.

Table III shows a list of the versions that were removed from the training set.

Training yielded expected results with the accuracy going down from 86.77% to 79.88% along with an increase in loss from 52% to 103%. This indicates that although the model is achieving almost 80% accuracy, it is also making some significant errors in its predictions, along with signs of overfitting.

The first test was done using the buffalo RGB set on the model-2 CNN. The results suggest a decrease in accuracy loss due to missing firmware, this decrease in accuracy was smaller than expected.

We decided to see if the loss could be brought down by increasing the depth and the best candidate for this would be model-4 due to its consistently high performance across all the vendors. After testing, model-4 seems to have traded some accuracy for better loss with the end results being [80.37% accuracy 89.28% loss] after 20 epocs.

Using a deeper model (model-4) results in a decrease in validation loss, overfitting is still an issue at this point. The increase in depth allowed the model to learn more complex features thereby improving the model's capability of predicting unencountered firmware signified by the reduction in validation loss. The accuracy on the other hand could be increased by calibrating the hyperparameters and number of epocs, while creating more augmented images for better representation in the dataset should improve the overfitting issue.

VII. CONCLUSION

A. Concluding Remarks

An IoT cyberattack can cause massive damage to the connected devices and harm to their owners. Thus, router firmware vulnerability detection is recently becoming an emerging issue in this domain. As a result, we provided an efficient and precise detection tool that is a necessity to this domain. The firmware dataset collection

was manually collected and made it available to the community. This paper investigated the effectiveness of analyzing home router firmware by employing the use of convolutional neural networks (CNN) and the use of computer vision techniques. We tested a set of different combinations of the filtered/normal datasets with the CNN training model. The highest accuracy filter was determined to be the HOG filter with an average accuracy of 85.81% across all tests and models with results as high as 97.94% when used with the right CNN model.

B. Limitations

During the data collection phase, no premade firmware datasets for home routers could be found online, resulting in grueling manual collection. Given more human and financial resources, a much larger dataset could be collected spanning a wider variety of router models. Due to the smaller dataset, some router models had to be removed as the preliminary tests had determined that underrepresented models instituted unnecessary variability and decreases in accuracy. Another limitation came from having less computational power. Having more computers would have allowed for faster progress during testing, and ultimately freed time for more filters to be tested.

In the end, the experimental results showed that filtering an image into one with more distinct and clear features results in higher accuracy when used on deeper CNN models.

C. Future Work

In the future, we plan to test more models and filters on the datasets. The next model to be tested is the gray level cooccurrence matrix, and is used in displaying the combinations of pixel values in various directions while outputting an image with smaller dimensions and much less noise data than the ones used in the previous tests. We also plan to use supervised learning models such as the kNN algorithm in order to find out if classifying all manufacturers at once is a viable option.

Moreover, we will use a set of highly optimized models for image classifications such as VGGNet16 and ResNet50.

REFERENCES

- [1] S. Alrabae, M. Debbabi, and L. Wang, "A survey of binary code fingerprinting approaches: taxonomy, methodologies, and features," *ACM Computing Surveys (CSUR)*, vol. 55, no. 1, pp. 1–41, 2022.
- [2] A. Alsuwaidi, A. Hassan, F. Alkhatri, H. Ali, Q. Mohammad, and S. Alrabae, "Security vulnerabilities detected in medical devices," in 2020 12th Annual Undergraduate Research Conference on Applied Computing (URC). IEEE, 2020, pp. 1–6.
- [3] K. Cheng, Q. Li, L. Wang, Q. Chen, Y. Zheng, L. Sun, and Z. Liang, "Dtaint: detecting the taint-style vulnerability in embedded device firmware," in 2018 48th Annual IEEE/IFIP International Conference on Dependable Systems and Networks (DSN). IEEE, 2018, pp. 430–441.
- [4] J. Gao, X. Yang, Y. Jiang, H. Song, K.-K. R. Choo, and J. Sun, "Semantic learning based cross-platform binary vulnerability search for iot devices," *IEEE Transactions on Industrial Informatics*, vol. 17, no. 2, pp. 971–979, 2019.
- [5] G. George and S. M. Thampi, "Combinatorial analysis for securing iot-assisted industry 4.0 applications from vulnerability-based attacks," *IEEE Transactions on Industrial Informatics*, 2020.
- [6] Y. Liu, J. Wang, J. Li, S. Niu, and H. Song, "Machine learning for the detection and identification of internet of things devices: A survey," *IEEE Internet of Things Journal*, vol. 9, no. 1, pp. 298–320, 2021.
- [7] —, "Class-incremental learning for wireless device identification in iot," *IEEE Internet of Things Journal*, vol. 8, no. 23, pp. 17 227–17 235, 2021.
- [8] Y. Liu, J. Wang, J. Li, S. Niu, L. Wu, and H. Song, "Zero-bias deep learning enabled quickest abnormal event detection in iot," *IEEE Internet of Things Journal*, 2021.

- [9] S. N. Matheu, J. L. Hernandez-Ramos, A. F. Skarmeta, and G. Baldini, "A survey of cybersecurity certification for the internet of things," *ACM Computing Surveys (CSUR)*, vol. 53, no. 6, pp. 1–36, 2020.
- [10] M. Poongodi, M. Hamdi, and H. Wang, "Image and audio caps: automated captioning of background sounds and images using deep learning," *Multimedia Systems*, pp. 1–9, 2022.
- [11] P. Srivastava, H. Peng, J. Li, H. Okhravi, H. Shrobe, and M. Payer, "Firmfuzz: Automated iot firmware introspection and analysis," in *Proceedings of the 2nd International ACM Workshop on Security and Privacy for the Internet-of-Things*, 2019, pp. 15–21.
- [12] J. Costoya, R. Flores, L. Gu, and F. Mercês, "Securing your home routers," *Web*, 2019.
- [13] B. Feng, A. Mera, and L. Lu, "P2IM: Scalable and hardware-independent firmware testing via automatic peripheral interface modeling," in *29th USENIX Security Symposium (USENIX Security 20)*, 2020, pp. 1237–1254.
- [14] M. Kim, D. Kim, E. Kim, S. Kim, Y. Jang, and Y. Kim, "Firmae: Towards large-scale emulation of iot firmware for dynamic analysis," in *Annual Computer Security Applications Conference*, 2020, pp. 733–745.
- [15] A. Qasem, P. Shirani, M. Debbabi, L. Wang, B. Lebel, and B. L. Agba, "Automatic vulnerability detection in embedded devices and firmware: Survey and layered taxonomies," *ACM Computing Surveys (CSUR)*, vol. 54, no. 2, pp. 1–42, 2021.
- [16] P. Shirani, L. Collard, B. L. Agba, B. Lebel, M. Debbabi, L. Wang, and A. Hanna, "B in a rm: Scalable and efficient detection of vulnerabilities in firmware images of intelligent electronic devices," in *Detection of Intrusions and Malware, and Vulnerability Assessment: 15th International Conference, DIMVA 2018, Saclay, France, June 28–29, 2018, Proceedings 15*. Springer, 2018, pp. 114–138.
- [17] J. Vadayath, M. Eckert, K. Zeng, N. Weideman, G. P. Menon, Y. Fratanonio, D. Balzarotti, A. Doupe, T. Bao, R. Wang et al, "Arbiter: Bridging the static and dynamic divide in vulnerability discovery on binary programs," in *31st USENIX Security Symposium (USENIX Security 22)*, 2022, pp. 413–430.
- [18] H. Wen, Z. Lin, and Y. Zhang, "Firmxray: Detecting bluetooth link layer vulnerabilities from bare-metal firmware," in *Proceedings of the 2020 ACM SIGSAC Conference on Computer and Communications Security*, 2020, pp. 167–180.
- [19] Y. Zheng, A. Davanian, H. Yin, C. Song, H. Zhu, and L. Sun, "FIRM-AFL: High-Throughput greybox fuzzing of IoT firmware via augmented process emulation," in *28th USENIX Security Symposium (USENIX Security 19)*, 2019, pp. 1099–1114.
- [20] L. Chen, Y. Wang, Q. Cai, Y. Zhan, H. Hu, J. Linghu, Q. Hou, C. Zhang, H. Duan, and Z. Xue, "Sharing more and checking less: Leveraging common input keywords to detect bugs in embedded systems," in *30th USENIX Security Symposium (USENIX Security 21)*, 2021.
- [21] N. Redini, A. Machiry, R. Wang, C. Spensky, A. Continella, Y. Shoshitaishvili, C. Kruegel, and G. Vigna, "Karonte: Detecting insecure multi-binary interactions in embedded firmware," in *2020 IEEE Symposium on Security and Privacy (SP)*. IEEE, 2020, pp. 1544–1561.
- [22] A. Pinhero, M. Anupama, P. Vinod, C. A. Visaggio, N. Aneesh, S. Abhijith, and S. AnanthaKrishnan, "Malware detection employed by visualization and deep neural network," *Computers Security*, p. 102247, 2021.
- [23] Z. Li, F. Liu, W. Yang, S. Peng, and J. Zhou, "A survey of convolutional neural networks: analysis, applications, and prospects," *IEEE Transactions on Neural Networks and Learning Systems*, 2021.
- [24] A. Mvoulana, R. Kachouri, and M. Akil, "Fine-tuning convolutional neural networks: a comprehensive guide and benchmark analysis for glaucoma screening," in *2020 25th International Conference on Pattern Recognition (ICPR)*. IEEE, 2021, pp. 6120–6127.
- [25] J. Brownlee, "How do convolutional layers work in deep learning neural networks?" *Machine Learning Mastery*, 2020.
- [26] N. Kawasaki, "Parametric study of thermal and chemical nonequilibrium nozzle flow," M.S. thesis, Dept. Electron. Eng., Osaka Univ., Osaka, Japan, 1993.
- [27] A. Harrison, private communication, May 1995.
- [28] B. Smith, "An approach to graphs of linear forms," unpublished.
- [29] A. Brahms, "Representation error for real numbers in binary computer arithmetic," *IEEE Computer Group Repository*, Paper R-67-85.
- [30] G.-B. Huang and H. A. Babri, "Upper bounds on the number of hidden neurons in feedforward networks with arbitrary bounded nonlinear activation functions," *IEEE transactions on neural networks*, vol. 9, no. 1, pp. 224–229, 1998.
- [31] F. Agostinelli, M. Hoffman, P. Sadowski, and P. Baldi, "Learning activation functions to improve deep neural networks," *arXiv preprint arXiv:1412.6830*, 2014.
- [32] R. Fardel, M. Nagel, F. Nuesch, T. Lippert, and A. Wokaun, "Fabrication of organic light emitting diode pixels by laser-assisted forward transfer," *Appl. Phys. Lett.*, vol. 91, no. 6, Aug. 2007, Art. no. 061103.
- [33] J. Zhang and N. Tansu, "Optical gain and laser characteristics of InGaN quantum wells on ternary InGa_N substrates," *IEEE Photon. J.*, vol. 5, no. 2, Apr. 2013, Art. no. 2600111.
- [34] S. Azodolmolky et al., Experimental demonstration of an impairment aware network planning and operation tool for transparent/translucent optical networks," *J. Lightw. Technol.*, vol. 29, no. 4, pp. 439–448, Sep. 2011.
- [35] J. Brownlee, "A gentle introduction to the rectified linear unit (relu)," *Machine learning mastery*, vol. 6, 2019.